

EXECUTIVE DOSSIER: DATA INTEGRITY & CLINICAL VALIDATION

Project: Health Performance Systems Optimization (5,000 Records)

Lead Analyst: Ayandokun Shadrach

Organization: DataVerse Africa / NovaCare Health Systems

Publication Date: May 8, 2026

Document Classification: Internal Operational Intelligence (Highly Confidential)

TABLE OF CONTENTS

1. EXECUTIVE SUMMARY

- 1.1 The Challenge: Data Entropy in High-Stakes Environments
- 1.2 The Mandate: Establishing the "Single Source of Truth" (SSOT)
- 1.3 Methodological Rigor
- 1.4 Strategic Impact

2. SYSTEMIC AUDIT & ROOT CAUSE ANALYSIS

- 2.1 The Readmission Paradox: Integrity vs. Administrative Lag
- 2.2 Demographic Structural Mismatch: The "Phantom" Pediatric Inflation
- 2.3 Logical Impossibilities: Medical-Legal Vulnerability
- 2.4 Operational Contradictions: Workflow Integrity Breach

3. DATA TRANSFORMATION PROTOCOL (ETL ARCHITECTURE)

- 3.1 Temporal Re-Engineering: The Readmission Logic
- 3.2 Clinical Logic Hierarchy: Establishing the "Source of Truth"
- 3.3 Demographic & Identity Reconciliation
 - 3.3.1 *Ward Purity & KPI Isolation*
 - 3.3.2 *Identity Normalization (Age & Gender)*

4. FINANCIAL INTEGRITY & STATISTICAL IMPUTATION

- 4.1 The Challenge of Missing Fiscal Data
- 4.2 The Imputation Strategy: Why Median over Mean?
- 4.3 Impact on Downstream Analytics

5. CLINICAL TAXONOMY MAPPING

- 5.1 Bridging the Semantic Gap
- 5.2 Operational Contextualization of ICD-10 & CPT Codes
- 5.3 Enhanced Business Intelligence (BI) Usability
- 5.4 Scalability of the Taxonomy Layer

6. QUALITY ASSURANCE SCORECARD & VALIDATION

- 6.1 The Zero-Tolerance Framework
- 6.2 Dimension 1: Logical Consistency
- 6.3 Dimension 2: Temporal Validity
- 6.4 Dimension 3: Categorical Accuracy (KPI Purity)
- 6.5 Stakeholder Readiness & Final Certification

7. CONCLUSION & STRATEGIC RECOMMENDATIONS

- **7.1** Final Assessment of Data Integrity
- **7.2** Strategic Recommendations for NovaCare Systems
- **7.3** Closing Statement

8. TECHNICAL APPENDIX: FORMULA REPOSITORY

- **8.1** Readmission Standardization Logic
- **8.2** Pediatric KPI Protection Formulas
- **8.3** Cost Normalization & Imputation Scripts
- **8.4** Temporal Metrics (LOS & ED Wait Time)

1. EXECUTIVE SUMMARY

1.1 The Challenge: Data Entropy in High-Stakes Environments In the modern healthcare landscape, "Data Entropy"—the natural decay of information quality due to fragmented entry points and system siloes—often obscures operational truth. At NovaCare, this entropy manifested as systemic contradictions between administrative tagging and actual clinical events. When administrative data (like billing flags) disagrees with clinical reality (like patient discharge status), the organization suffers from "Information Asymmetry," leading to poor resource allocation, lost revenue, and compromised patient safety metrics.

1.2 The Mandate: Establishing the "Single Source of Truth" (SSOT) This dossier documents a rigorous, multi-stage sanitization and clinical validation initiative involving 5,000 unique hospital encounter records. The primary objective was not merely to "clean" data, but to perform **Clinical Data Reconciliation**. By resolving deep-seated contradictions—such as "Readmission Paradoxes" and "Clinical Impossibilities"—we have engineered a robust analytical foundation. This project establishes a "Single Source of Truth" (SSOT) for NovaCare's performance analytics, ensuring that executive decisions are based on validated evidence rather than legacy system errors.

1.3 Methodological Rigor The transformation process utilized advanced ETL (Extract, Transform, Load) architectures, moving beyond simple spreadsheet edits. We implemented:

- **Temporal Re-Engineering:** Reconstructing patient timelines from raw timestamps to verify readmission accuracy.
- **Cross-Domain Validation:** Triangulating data between Finance (Cost), Pharmacy/Lab (Procedures), and Clinical Wards (Outcome/Disposition).
- **Statistical Imputation:** Utilizing median-based distribution modeling to address missing financial data without introducing the bias inherent in traditional mean-averaging.

1.4 Strategic Impact The result of this initiative is a high-fidelity dataset that serves as a "Clean Pipe" for downstream Business Intelligence (BI). With this foundation, NovaCare can now achieve:

- **Precision Revenue Cycle Management:** Identifying leakage with 100% data confidence.
- **Operational Transparency:** Accurate tracking of Departmental Efficiency (LOS) and Patient Throughput.
- **Audit-Ready Compliance:** Establishing a defensible data trail that meets the highest standards of healthcare governance.

2. SYSTEMIC AUDIT & ROOT CAUSE ANALYSIS (EXPANDED)

2.1 The Readmission Paradox: Integrity vs. Administrative Lag The audit identified a 19% discrepancy rate between the `Outcome` field and the `Readmission_Flag`. In clinical informatics, this is often the result of "Status Fragmentation," where a patient is readmitted via the Emergency Department (ED), but the administrative flag fails to trigger because the previous encounter was not yet "closed" in the billing system.

- **The Risk:** Relying on the raw flag would lead to a significant underreporting of hospital-acquired complications and readmission penalties, skewing the hospital's Quality of Care metrics.
- **Root Cause:** Lack of real-time synchronization between the Clinical EMR and the Revenue Cycle Management (RCM) software.

2.2 Demographic Structural Mismatch: The "Phantom" Pediatric Inflation The discovery that 92.6% of the Pediatrics ward volume consisted of adults (Ages 18–99) highlights a critical "Boarding" crisis. In many high-volume facilities, adult medical-surgical wards reach 100% capacity, forcing the "overflow" of geriatric or adult patients into specialized pediatric beds.

- **The Risk:** Pediatric Length-of-Stay (LOS) and "Cost per Case" KPIs become meaningless. An 80-year-old with a UTI has a vastly different clinical recovery path than a 5-year-old; mixing these datasets creates a "phantom" inflation of pediatric stay durations, making the department appear inefficient when it is actually being used as an overflow valve.
- **Root Cause:** Acute bed shortages in Internal Medicine and a lack of automated ward-validation rules at the point of admission.

2.3 Logical Impossibilities: Medical-Legal Vulnerability The mapping of 28 "Expired" outcomes to a "Home" or "Rehab" disposition is the highest-tier data failure. In a clinical setting, an "Expired" status must trigger a specific "Mortuary" or "Funeral Home" disposition.

- **The Risk:** Beyond data quality, this represents a severe medical-legal risk. If a deceased patient is recorded as "Discharged to Home," it creates a fraudulent audit trail that could lead to regulatory fines, insurance recoupment, and legal liability.
- **Root Cause:** Manual data entry overrides where the user selects the top-listed disposition (usually "Home") without verifying the clinical outcome.

2.4 Operational Contradictions: Workflow Integrity Breach The classification of 286 Emergency Department (ED) cases as "Elective" represents a fundamental breach of clinical definitions. By definition, an ED admission is unplanned/acute, whereas "Elective" refers to scheduled procedures (e.g., a planned knee replacement).

- **The Risk:** This error masks the true burden of "Unscheduled Care." If ED cases are labeled elective, hospital leadership cannot accurately calculate the "Surge Capacity" needed for real emergencies, leading to staffing shortages and increased wait times.
- **Root Cause:** A default system setting in the admission module that fails to update the "Admission Type" based on the originating department.

3. DATA TRANSFORMATION PROTOCOL (ETL ARCHITECTURE)

3.1 Temporal Re-Engineering: The Readmission Logic

In clinical analytics, the "Readmission Rate" is a critical proxy for quality of care. Because the raw dataset contained inconsistent flags, we bypassed existing administrative labels in favor of a "First-Principles" chronological reconstruction.

- **Methodology (Partitioned Sorting):** Data was partitioned by `Patient_ID` and sorted by `Admission_Datetime`. This allowed us to view each patient's journey as a linear sequence of events rather than isolated rows.
- **The Calculation (Days_Since_Last_Exit):** We engineered a `Previous_Discharge_Date` variable using a `MAXIFS` array formula. This formula identifies the most recent discharge date that occurred *prior* to the current admission, ensuring that overlapping or incorrectly entered dates did not create negative "days-between" values.
- **The Industry Standard:** By calculating the delta between the current admission and the previous discharge, we validated a new "Ground-Truth" flag. Any patient returning within ≤ 30 days was marked as a clinical readmission.

3.2 Clinical Logic Hierarchy: Establishing the "Source of Truth"

A major challenge in healthcare data is "Source Conflict," where different modules (e.g., Clinical vs. Billing) provide contradictory statuses. We implemented a strict hierarchy of truth to resolve these conflicts:

- **The Biological Integrity Rule:** Clinical status is the ultimate authority. If a patient's `Outcome` is recorded as "Expired," the `Discharge_Disposition` must reflect this. We force-corrected any "Home" or "Rehab" dispositions to "Expired" to prevent legal and audit discrepancies.
- **Cross-Field Synchronization:** To ensure the Power BI dashboard provides a unified view, we synchronized the `Outcome` field with our new Readmission Flag. If a patient met the 30-day criteria, their outcome was standardized to "Readmitted," ensuring that "Recovered" labels did not mask the reality of a bounce-back encounter.

3.3 Demographic & Identity Reconciliation

To ensure the dataset is "Analytic-Ready," we addressed structural mismatches in how patient profiles were recorded.

- **Ward Purity & KPI Isolation:** To protect the integrity of Pediatric Average Length of Stay (ALOS), we isolated the 463 adults "boarded" in the Pediatric ward. By moving

them to a virtual "**Misclassified**" ward, we removed the statistical noise they created in pediatric metrics while still retaining their financial contribution to the total hospital revenue cycle.

- **Identity Normalization:** Patient profiles often suffer from "Data Drifting," where a single patient may have slightly different ages or genders recorded across different years.
 - **Age:** We applied the **Median Age** per patient to neutralize typos (e.g., a patient recorded as 45, 46, and 450).
 - **Gender:** We utilized a **Mode-Based Gender** (Highest Frequency) approach. In the rare event of a tie, the system defaults to the earliest recorded entry, ensuring that each `Patient_ID` maps to a single, consistent demographic profile.

4. FINANCIAL INTEGRITY & STATISTICAL IMPUTATION

4.1 The Challenge of Missing Fiscal Data In hospital revenue cycle management, missing cost data is not merely a technical "null" value; it represents a break in the financial audit trail. For NovaCare, 120 encounters lacked recorded costs. Leaving these as null would have artificially deflated total revenue projections and skewed "Cost per Department" analyses.

4.2 The Imputation Strategy: Why Median over Mean? We applied a **Median Imputation (\$1,551.65)** to resolve these gaps. This was a deliberate choice based on the statistical distribution of healthcare billing data.

- **The Right-Skewness of Healthcare Costs:** Hospital billing data is almost never "normally distributed." While most encounters (like routine consultations or minor labs) fall into a lower price bracket, a small number of high-intensity surgical procedures or long-term ICU stays create "fat tails" on the high end of the scale.
- **Neutralizing Outlier Bias:** Utilizing a mathematical **Mean** (average) would have been mathematically irresponsible. The high-cost outliers would have pulled the average upward, creating an "Upward Bias." This would have resulted in an overestimation of the typical cost for the 120 missing encounters.
- **The "Typical Patient" Proxy:** The **Median** represents the true 50th percentile. By using \$1,551.65, we ensured that our imputed values reflected the most frequent, "typical" encounter cost, thereby protecting the financial model from artificial inflation.

4.3 Impact on Downstream Analytics By stabilizing the financial column, we enabled reliable **Cost-to-Outcome** modeling. Stakeholders can now view the "Total Value at Risk" and "Revenue Leakage" KPIs with the confidence that the data has been normalized against extreme statistical noise, ensuring that financial forecasting for Q3 and Q4 remains grounded in clinical reality.

5. CLINICAL TAXONOMY MAPPING

5.1 Bridging the Semantic Gap In healthcare informatics, a significant barrier to executive decision-making is the "Semantic Gap"—the disconnect between cryptic billing codes (ICD-10/CPT) and clinical reality. Raw data containing codes like "E11" or "99284" is often unintelligible to non-clinical stakeholders, leading to under-utilized dashboards. To resolve this, we implemented a comprehensive **Taxonomy Translation Layer**.

5.2 Operational Contextualization We moved beyond simple labeling by mapping each technical code to its specific **Operational Context**. This allows leadership to view the hospital not as a collection of codes, but as a series of clinical service lines:

Technical Code	Clinical Translation	Operational Context
E11	Type 2 Diabetes Mellitus	Chronic Care Management: Enables tracking of long-term patient health trends and outpatient follow-up efficiency.
N39	Urinary Tract Infection (UTI)	High-Volume Acute Care: Facilitates analysis of rapid-turnover cases and antibiotic stewardship programs.
93000	Electrocardiogram (ECG)	Diagnostic Utilization: Measures the workload on ancillary services and cardiac diagnostic throughput.
99284	Emergency Dept Visit	Facility Throughput: Key for identifying Emergency Department "boarding" times and acute triage efficiency.

5.3 Enhanced Business Intelligence (BI) Usability By embedding these narratives directly into the data model via a centralized **Master Lookup Table**, we have enabled high-level "Drill-Down" capabilities in the Power BI reporting layer.

- **Stakeholder Impact:** An executive can now filter by "Chronic Care Management" to see the financial impact of Diabetes, without needing to know the specific ICD-10 codes involved.

- **Accuracy Guarantee:** This mapping was validated against standard WHO (World Health Organization) and AMA (American Medical Association) guidelines to ensure that the "Readable" names carry full clinical authority.

5.4 Scalability of the Taxonomy Layer This architecture is designed for scalability. As NovaCare expands its service lines (e.g., adding Oncology or Neurology), new codes can be integrated into the existing Lookup Table without requiring a restructure of the core data cleaning pipeline. This ensures that the "Single Source of Truth" remains adaptable to the hospital's evolving clinical landscape.

6. QUALITY ASSURANCE SCORECARD & VALIDATION

6.1 The Zero-Tolerance Framework

In high-stakes clinical environments, "mostly clean" data is a liability. To ensure the Health Performance Dataset met the rigorous standards required for hospital leadership, we moved beyond basic spot-checks to a **Zero-Tolerance Validation Framework**. This stage involved stress-testing the 5,000 records against three non-negotiable dimensions of data quality.

6.2 Dimension 1: Logical Consistency (The Clinical-Administrative Bridge)

- **The Standard:** Absolute alignment between clinical status and physical disposition.
- **The Validation:** We ran a cross-field audit ensuring that every patient marked as "Expired" in the clinical outcome was also marked as "Expired" in the discharge disposition.
- **The Outcome:** We achieved 100% compliance. By eliminating cases where deceased patients were erroneously recorded as "Discharged Home," we have mitigated significant medical-legal and financial audit risks for NovaCare.

6.3 Dimension 2: Temporal Validity (The Chronological Audit)

- **The Standard:** Zero instances of "Time-Travel" or negative durations.
- **The Validation:** We audited the `LOS_Hours` (Length of Stay) and `ED Wait Time` columns. In raw healthcare data, it is common to find discharge dates that occur before admission dates due to manual entry errors.
- **The Outcome:** Through our ETL pipeline, we identified and neutralized all temporal anomalies. Every encounter now follows a logical, linear path: *ED Arrival* $\rightarrow$ *Provider Seen* $\rightarrow$ *Admission* $\rightarrow$ *Discharge*. This ensures that departmental efficiency metrics are statistically defensible.

6.4 Dimension 3: Categorical Accuracy (KPI Purity)

- **The Standard:** Demographic alignment with departmental specialization.
- **The Validation:** We performed a "Ward-to-Age" reconciliation. Previously, the Pediatric ward was cluttered with adult geriatric data, skewing the department's performance profile.
- **The Outcome:** Pediatric performance metrics are now derived **exclusively from patients aged 0–17**. By reclassifying the 463 adult "overflow" patients, we have restored the integrity of the Pediatric Length of Stay (ALOS) KPI, allowing the department head to see the true efficiency of pediatric care for the first time.

6.5 Stakeholder Readiness & Final Certification

- **Metadata Enrichment:** Every ICD-10 and CPT code has been verified against the master lookup table.

- **Final Certification:** The dataset has been signed off as "**Production-Ready.**" It is now fully optimized for the Power BI reporting layer, ensuring that when a stakeholder filters by "Readmission Rate" or "Average Cost," the underlying data is a precise reflection of hospital operations.

7. CONCLUSION & STRATEGIC RECOMMENDATIONS

7.1 Final Assessment of Data Integrity

The comprehensive cleansing of the Health Performance Dataset marks a critical pivot from speculative reporting to evidence-based intelligence. By successfully mitigating the "Questionable" status of the raw source data, we have moved beyond simple error correction and into the realm of **Data Governance**. The resolution of systemic contradictions—specifically the reconciliation of clinical outcomes with administrative discharge logs—has transformed a fragmented dataset into a high-fidelity asset. We have developed a defensible, audit-ready framework that can withstand the scrutiny of both medical boards and financial auditors.

7.2 Strategic Recommendations for NovaCare Systems

- **Institutionalization of the "Clean_Data" Architecture:**

This architecture should serve as the mandated, authoritative foundation for all subsequent NovaCare operational dashboards. Any future predictive modeling, including patient-flow forecasting and machine learning-driven denial analytics, must ingest data through this validated pipeline to ensure model accuracy and clinical safety.

- **Implementation of Real-Time Validation Triggers:**

To prevent the re-emergence of "Data Entropy," we recommend implementing front-end validation rules within the Electronic Medical Record (EMR) system. Specifically, the system should prevent a "Home" discharge selection if the clinical outcome is marked as "Expired," and trigger an alert when an adult patient is assigned to a Pediatric bed.

- **Transition to Predictive Resource Allocation:**

With the establishment of the 30-day "Ground Truth" for readmissions, NovaCare is now positioned to move from descriptive analytics to **Prescriptive Analytics**. We recommend utilizing the standardized readmission data to identify "High-Risk Bounce-back" patient profiles, allowing for targeted post-discharge interventions that improve patient outcomes while reducing financial penalties.

- **Continuous Taxonomy Synchronization:**

The Clinical Taxonomy Mapping layer should be updated quarterly. As new clinical procedures are introduced, the mapping must ensure that technical CPT and ICD-10 codes remain transparent and actionable for non-clinical executive leadership.

7.3 Closing Statement

The success of this initiative demonstrates that the value of healthcare data is not found in its volume, but in its **validity**. By establishing this "Single Source of Truth," NovaCare is now

equipped to lead with precision, ensuring that every operational decision is backed by a clinical foundation that is as accurate as the care provided at the bedside.

8. TECHNICAL APPENDIX: FORMULA REPOSITORY

This appendix serves as the technical "Source of Truth" for the cleaning operations. These formulas were engineered to ensure that the ETL pipeline is repeatable, transparent, and mathematically sound.

8.1 Readmission Standardization Logic To resolve the conflict between existing outcome labels and actual patient behavior, we utilized a multi-step logic to identify "true" readmissions.

- **Previous Discharge Lookup:** `=IFERROR(MAXIFS(Discharge_Date_Range, Patient_ID_Range, Current_Patient_ID, Admission_Date_Range, "<"&Current_Admission_Date), "")`
 - *Purpose:* Identifies the most recent exit date for a specific patient prior to their current arrival.
- **Readmission Flag (30-Day Rule):** `=IF(AND(Days_Since_Last_Exit > 0, Days_Since_Last_Exit <= 30), 1, 0)`
- **Outcome Synchronization:** `=IF(Readmission_Flag = 1, "Readmitted", Original_Outcome)`
 - *Strategic Impact:* Corrects the 19% discrepancy rate identified during the audit.

8.2 Pediatric KPI Protection Formulas To prevent adult overflow from skewing pediatric performance metrics, we implemented a departmental reclassification script.

- **The Reclassification Rule:** `=IF(AND(Department = "Pediatrics", Age >= 18), "Misclassified", Department)`
 - *Purpose:* Dynamically shifts adult patients out of the pediatric denominator.
- **Clean Pediatric ALOS (Average Length of Stay):** `=AVERAGEIF(Department_Range, "Pediatrics", LOS_Hours_Range)`
 - *Strategic Impact:* Restores the integrity of pediatric efficiency reporting.

8.3 Cost Normalization & Imputation Scripts To address the 120 missing financial records without introducing the "Mean-Bias" common in skewed healthcare data.

- **Median Imputation Formula:** `=IF(ISBLANK(Cost_of_Encounter), 1551.65, Cost_of_Encounter)`
 - *Calculated Median:* \$1,551.65 (derived from the 50th percentile of all non-null encounters).
- **Total Value at Risk (DAX for Power BI):** `Total Value at Risk = SUMX(Filtered_Claims, [Paid_Amount] * [Denial_Probability])`
 - *Strategic Impact:* Protects the financial model from artificial inflation caused by high-cost surgical outliers.

8.4 Temporal Metrics (LOS & ED Wait Time) We converted raw timestamps into actionable "Throughput" metrics to identify facility bottlenecks.

- **Length of Stay (LOS) in Decimal Hours:** `=ROUND((Discharge_Datetime - Admission_Datetime) * 24, 2)`
 - *Purpose:* Provides high-resolution data for bed-turnover analysis.

- **ED Wait Time (Triage to Provider) in Minutes:** =ROUND((Provider_Seen_Time - ED_Arrival_Time) * 1440, 0)
 - *Purpose:* Standardizes the "Wait Time" KPI across all 5,000 records to identify the causes of the "Sunday Surge."
- **Readmission Interval:** =Admission_Datetime - Previous_Discharge_Date

Lead Analyst: Ayandokun Shadrach

Technical Verification: DataVerse Africa QA Protocol

Implementation Date: May 8, 2026